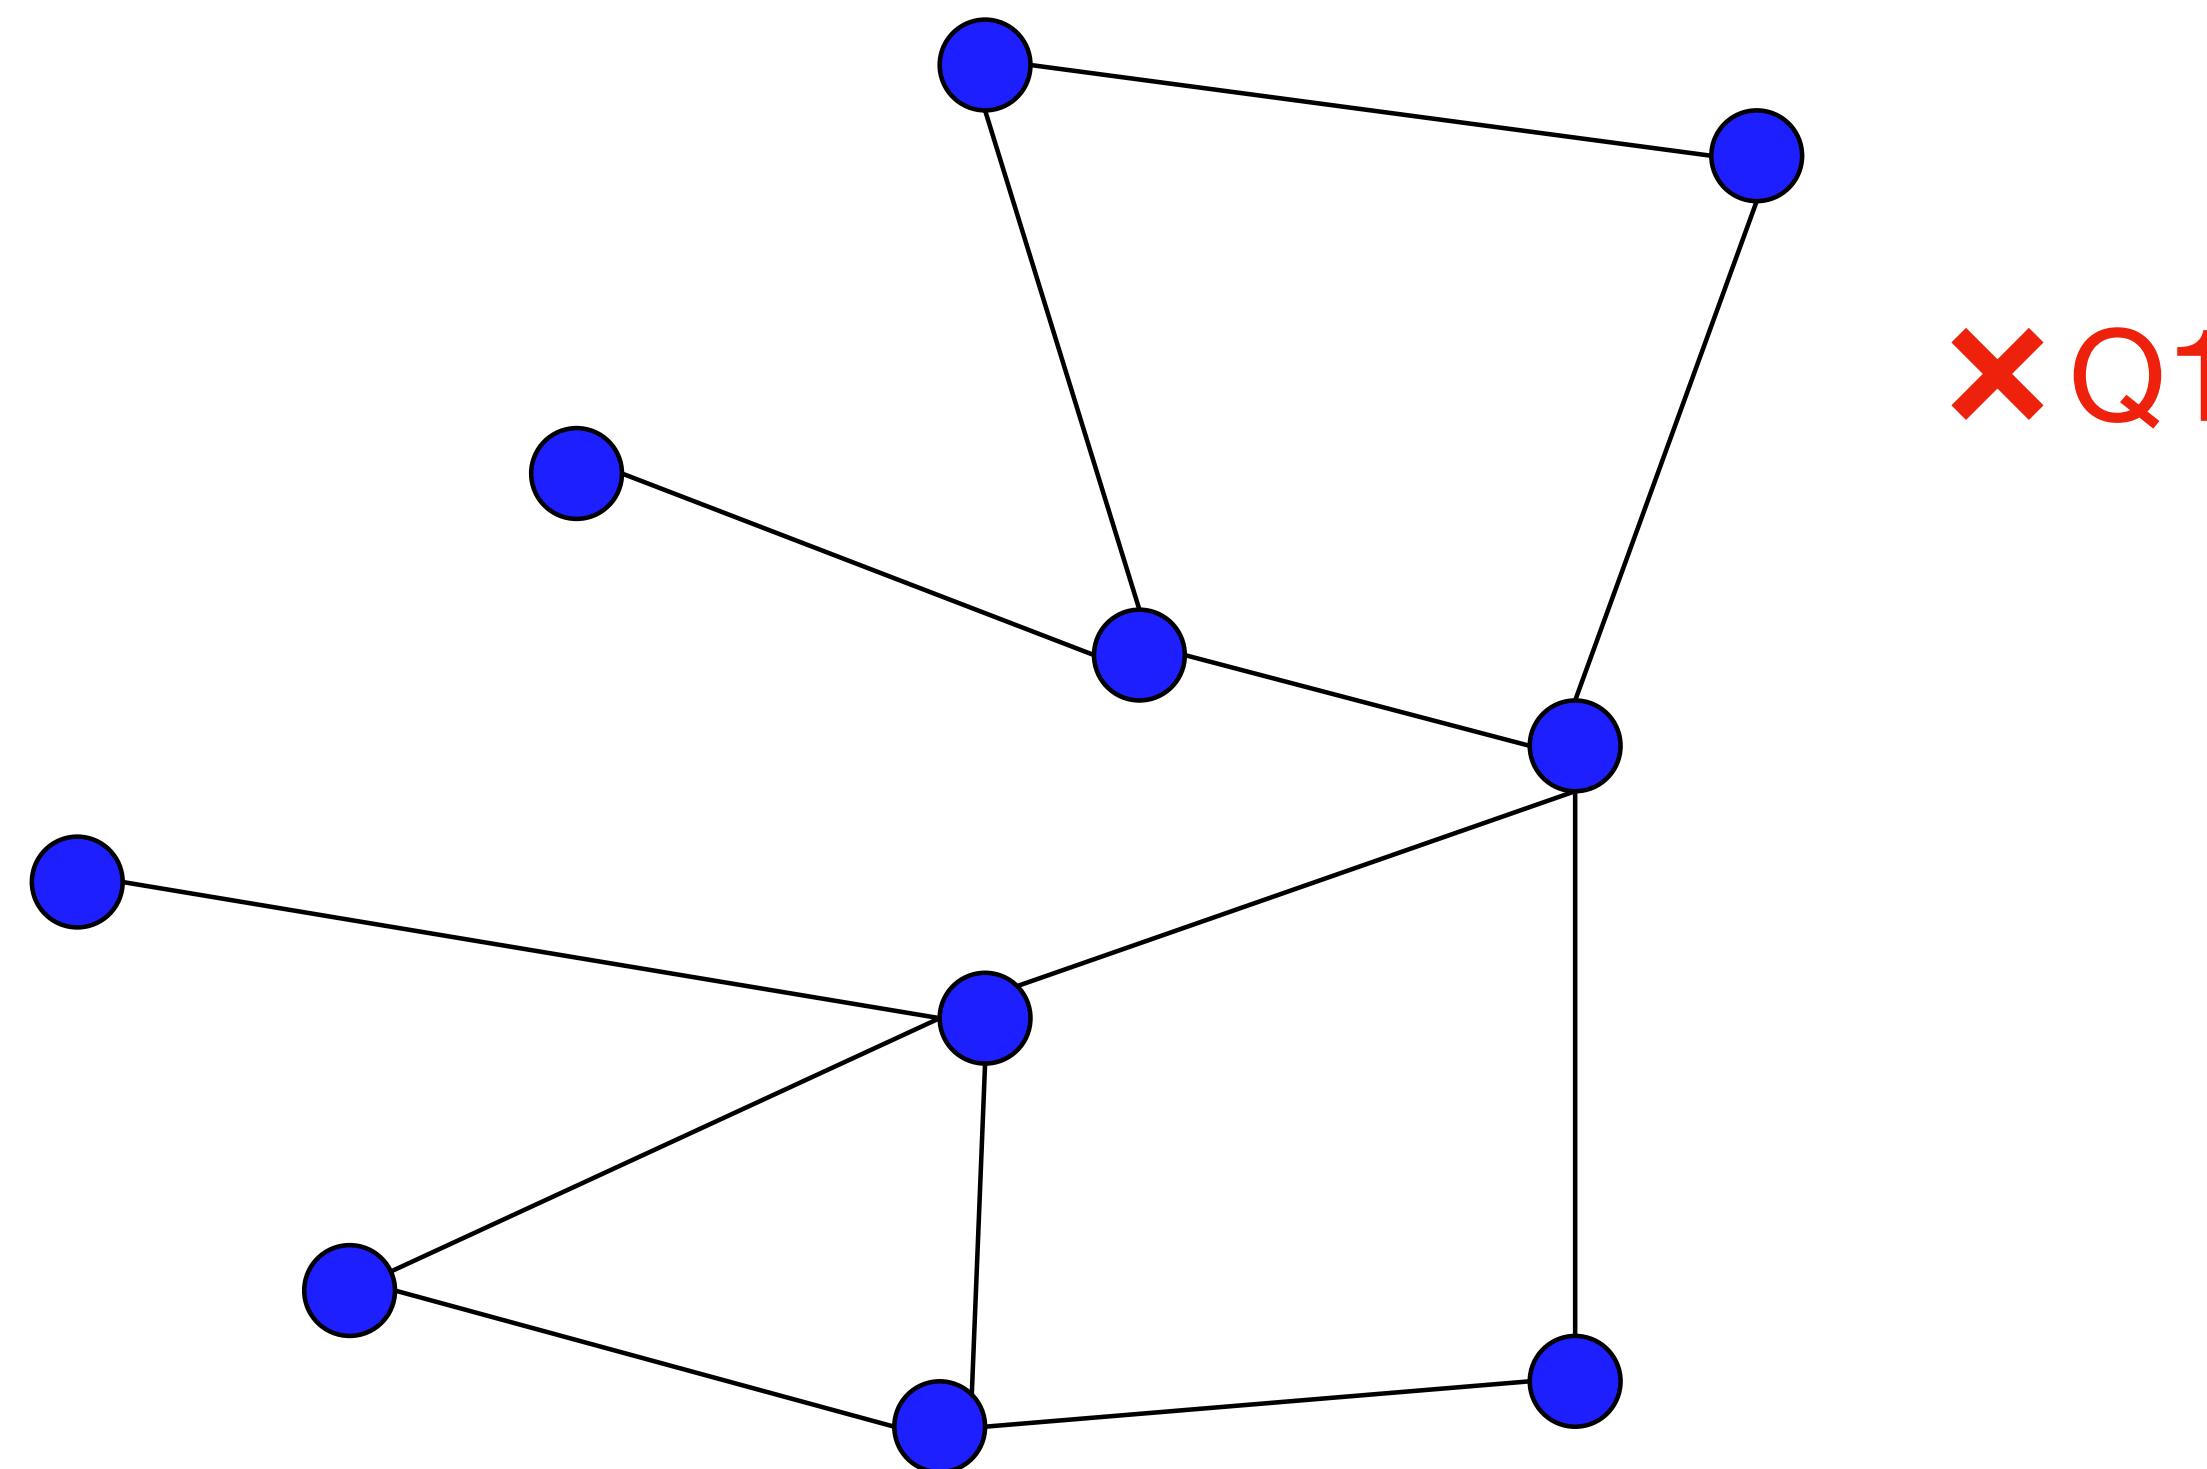


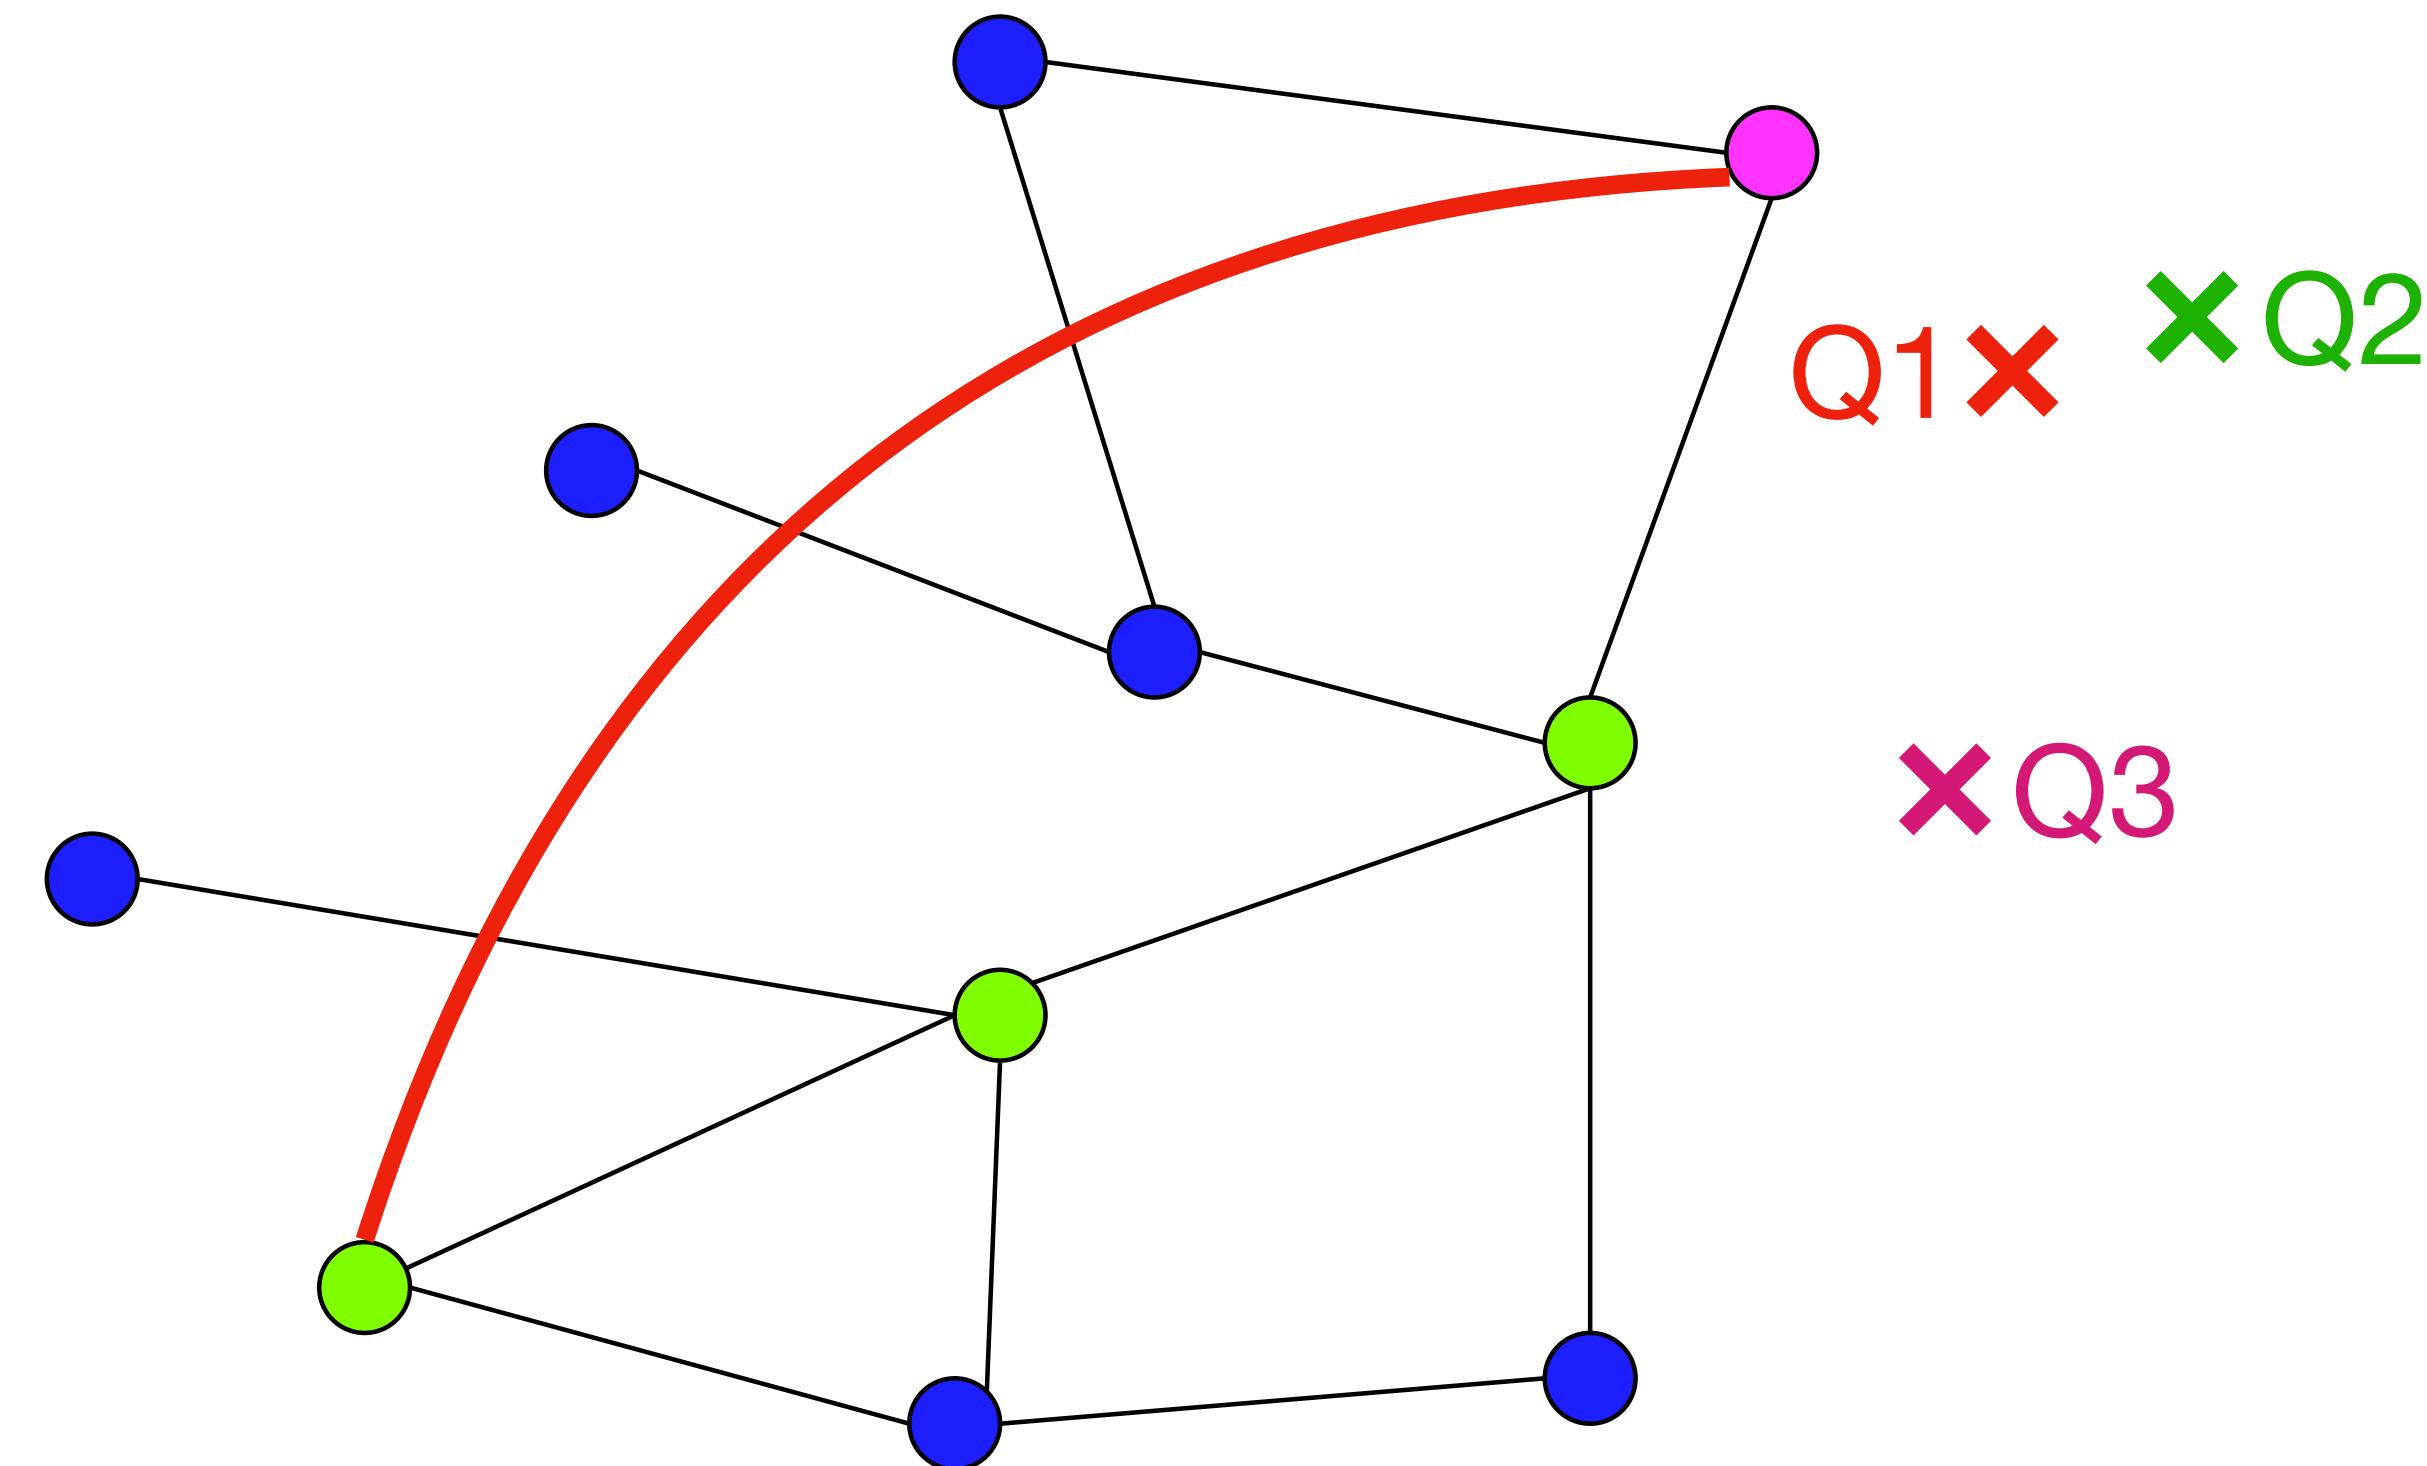
Project Catapult

...one month later

Graph databases



Catapults to the rescue



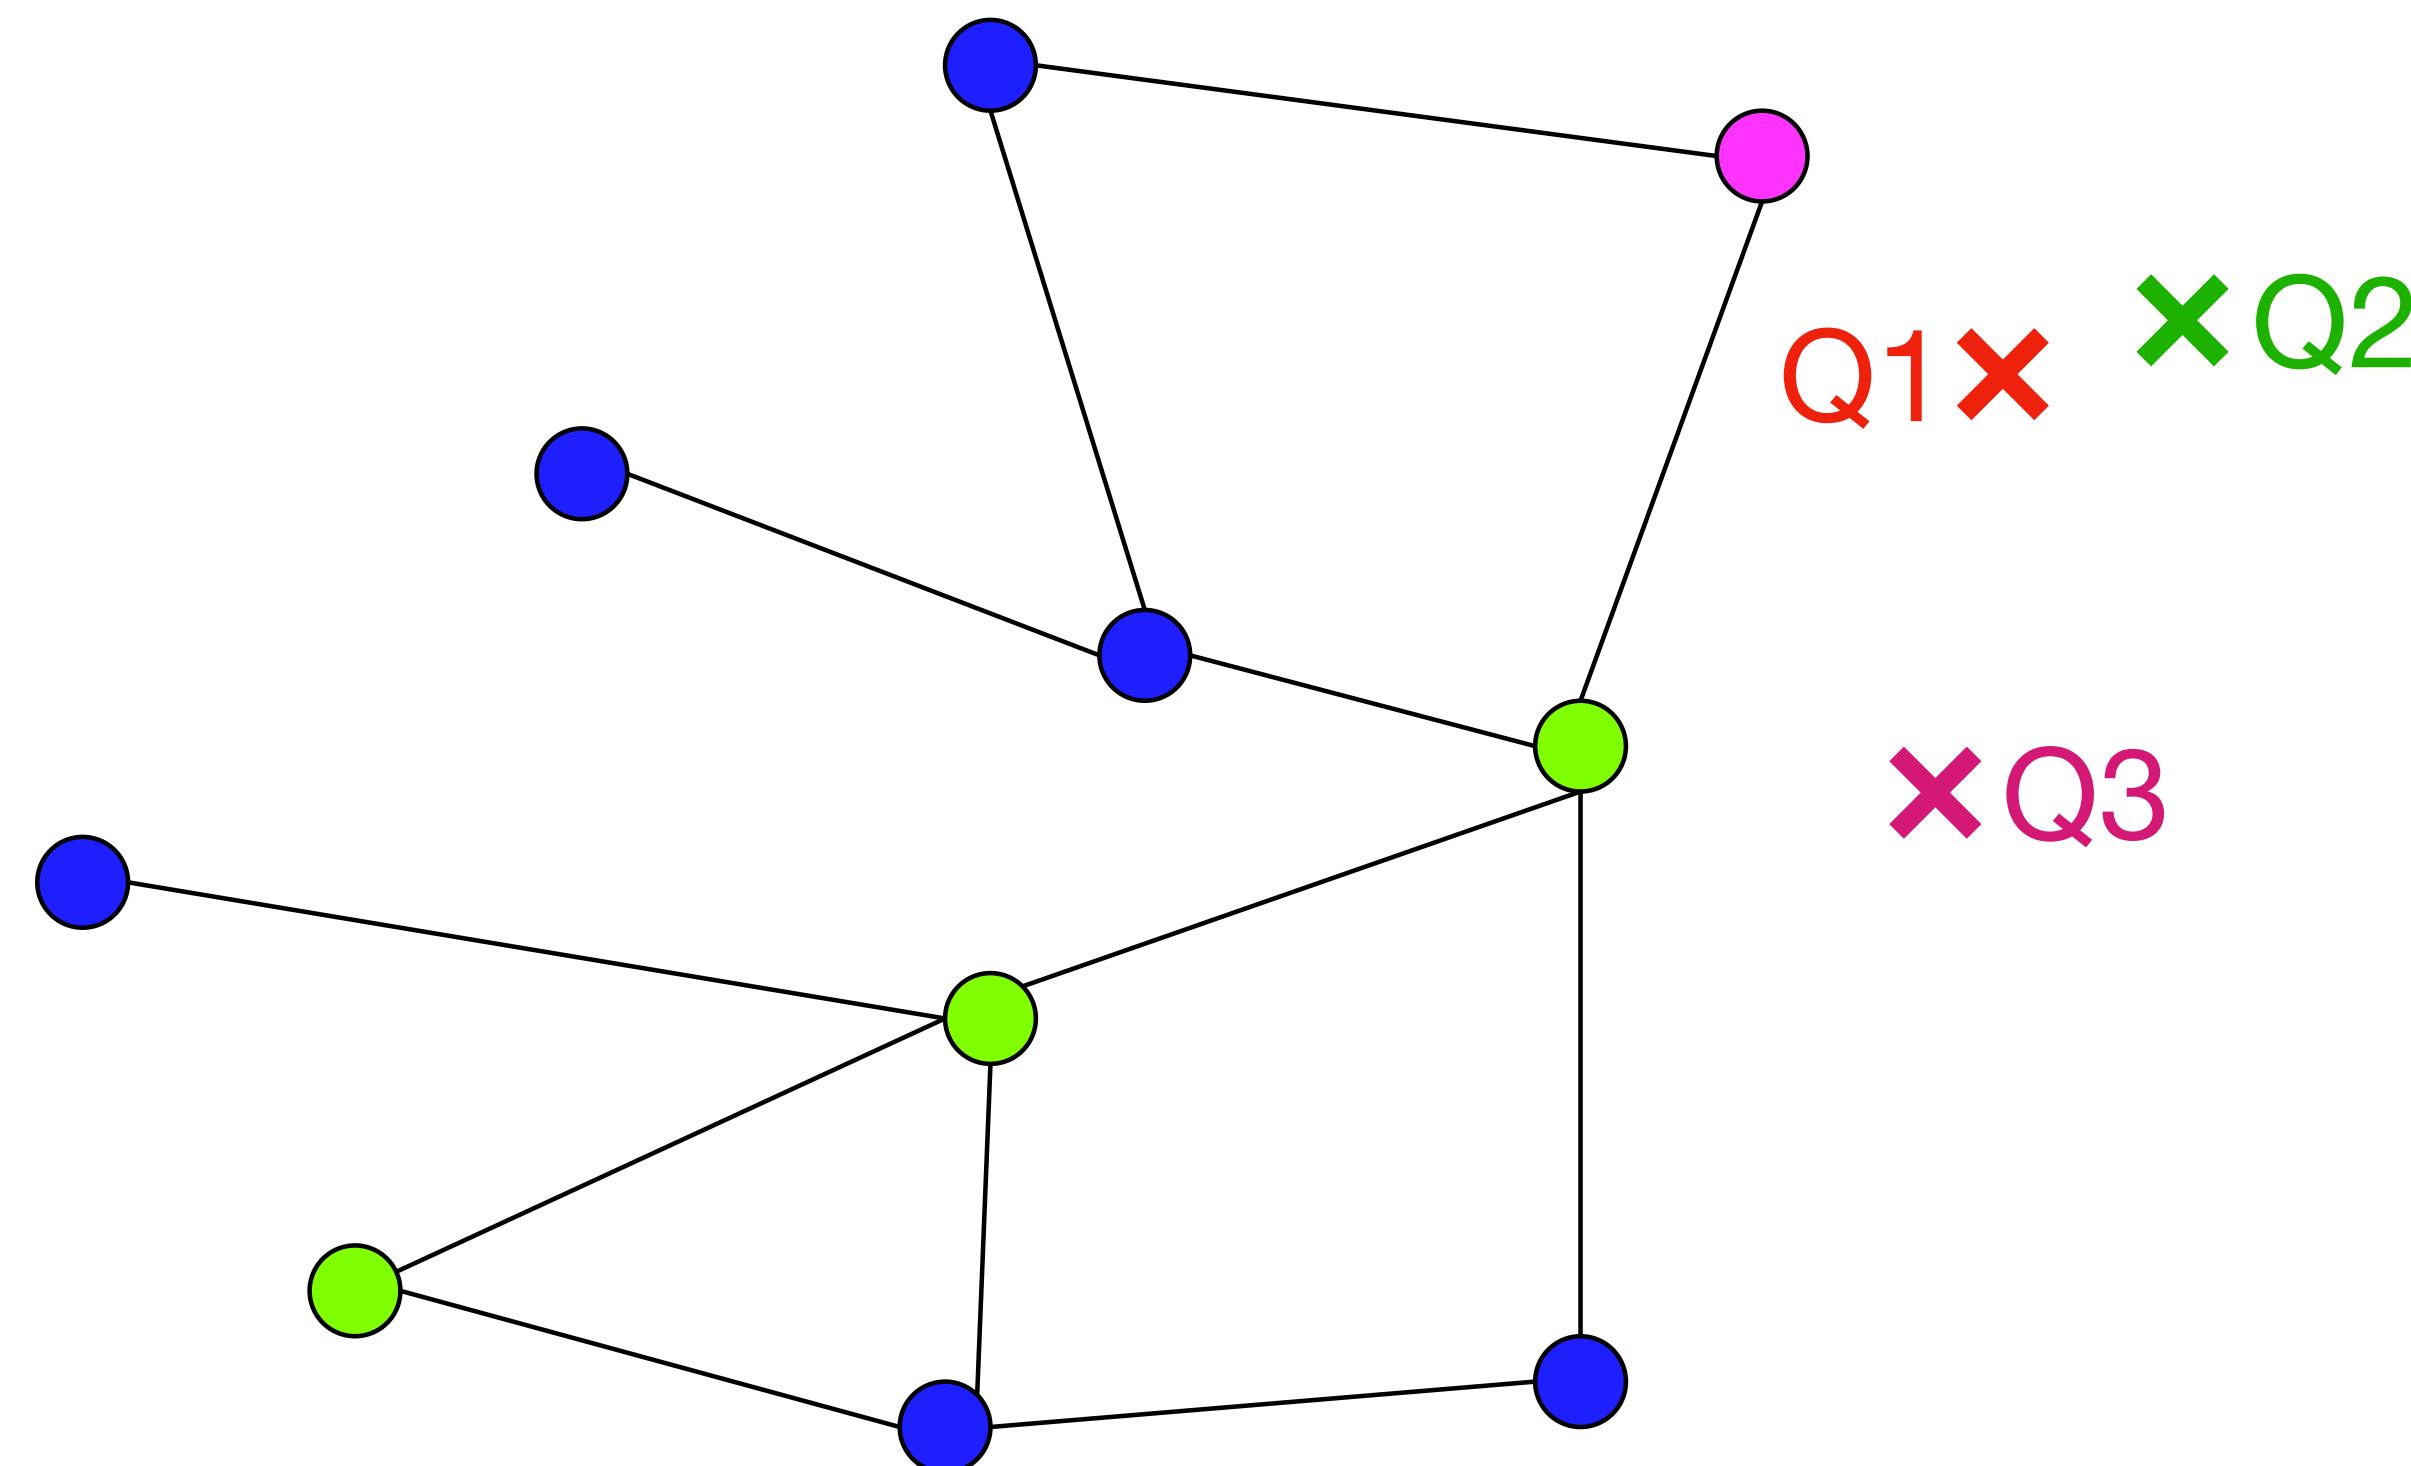
What do we have now?

What do we have?

- Search engine operational, including catapults
- Some things still WIP
- Works on toy data / large-scale experiments coming

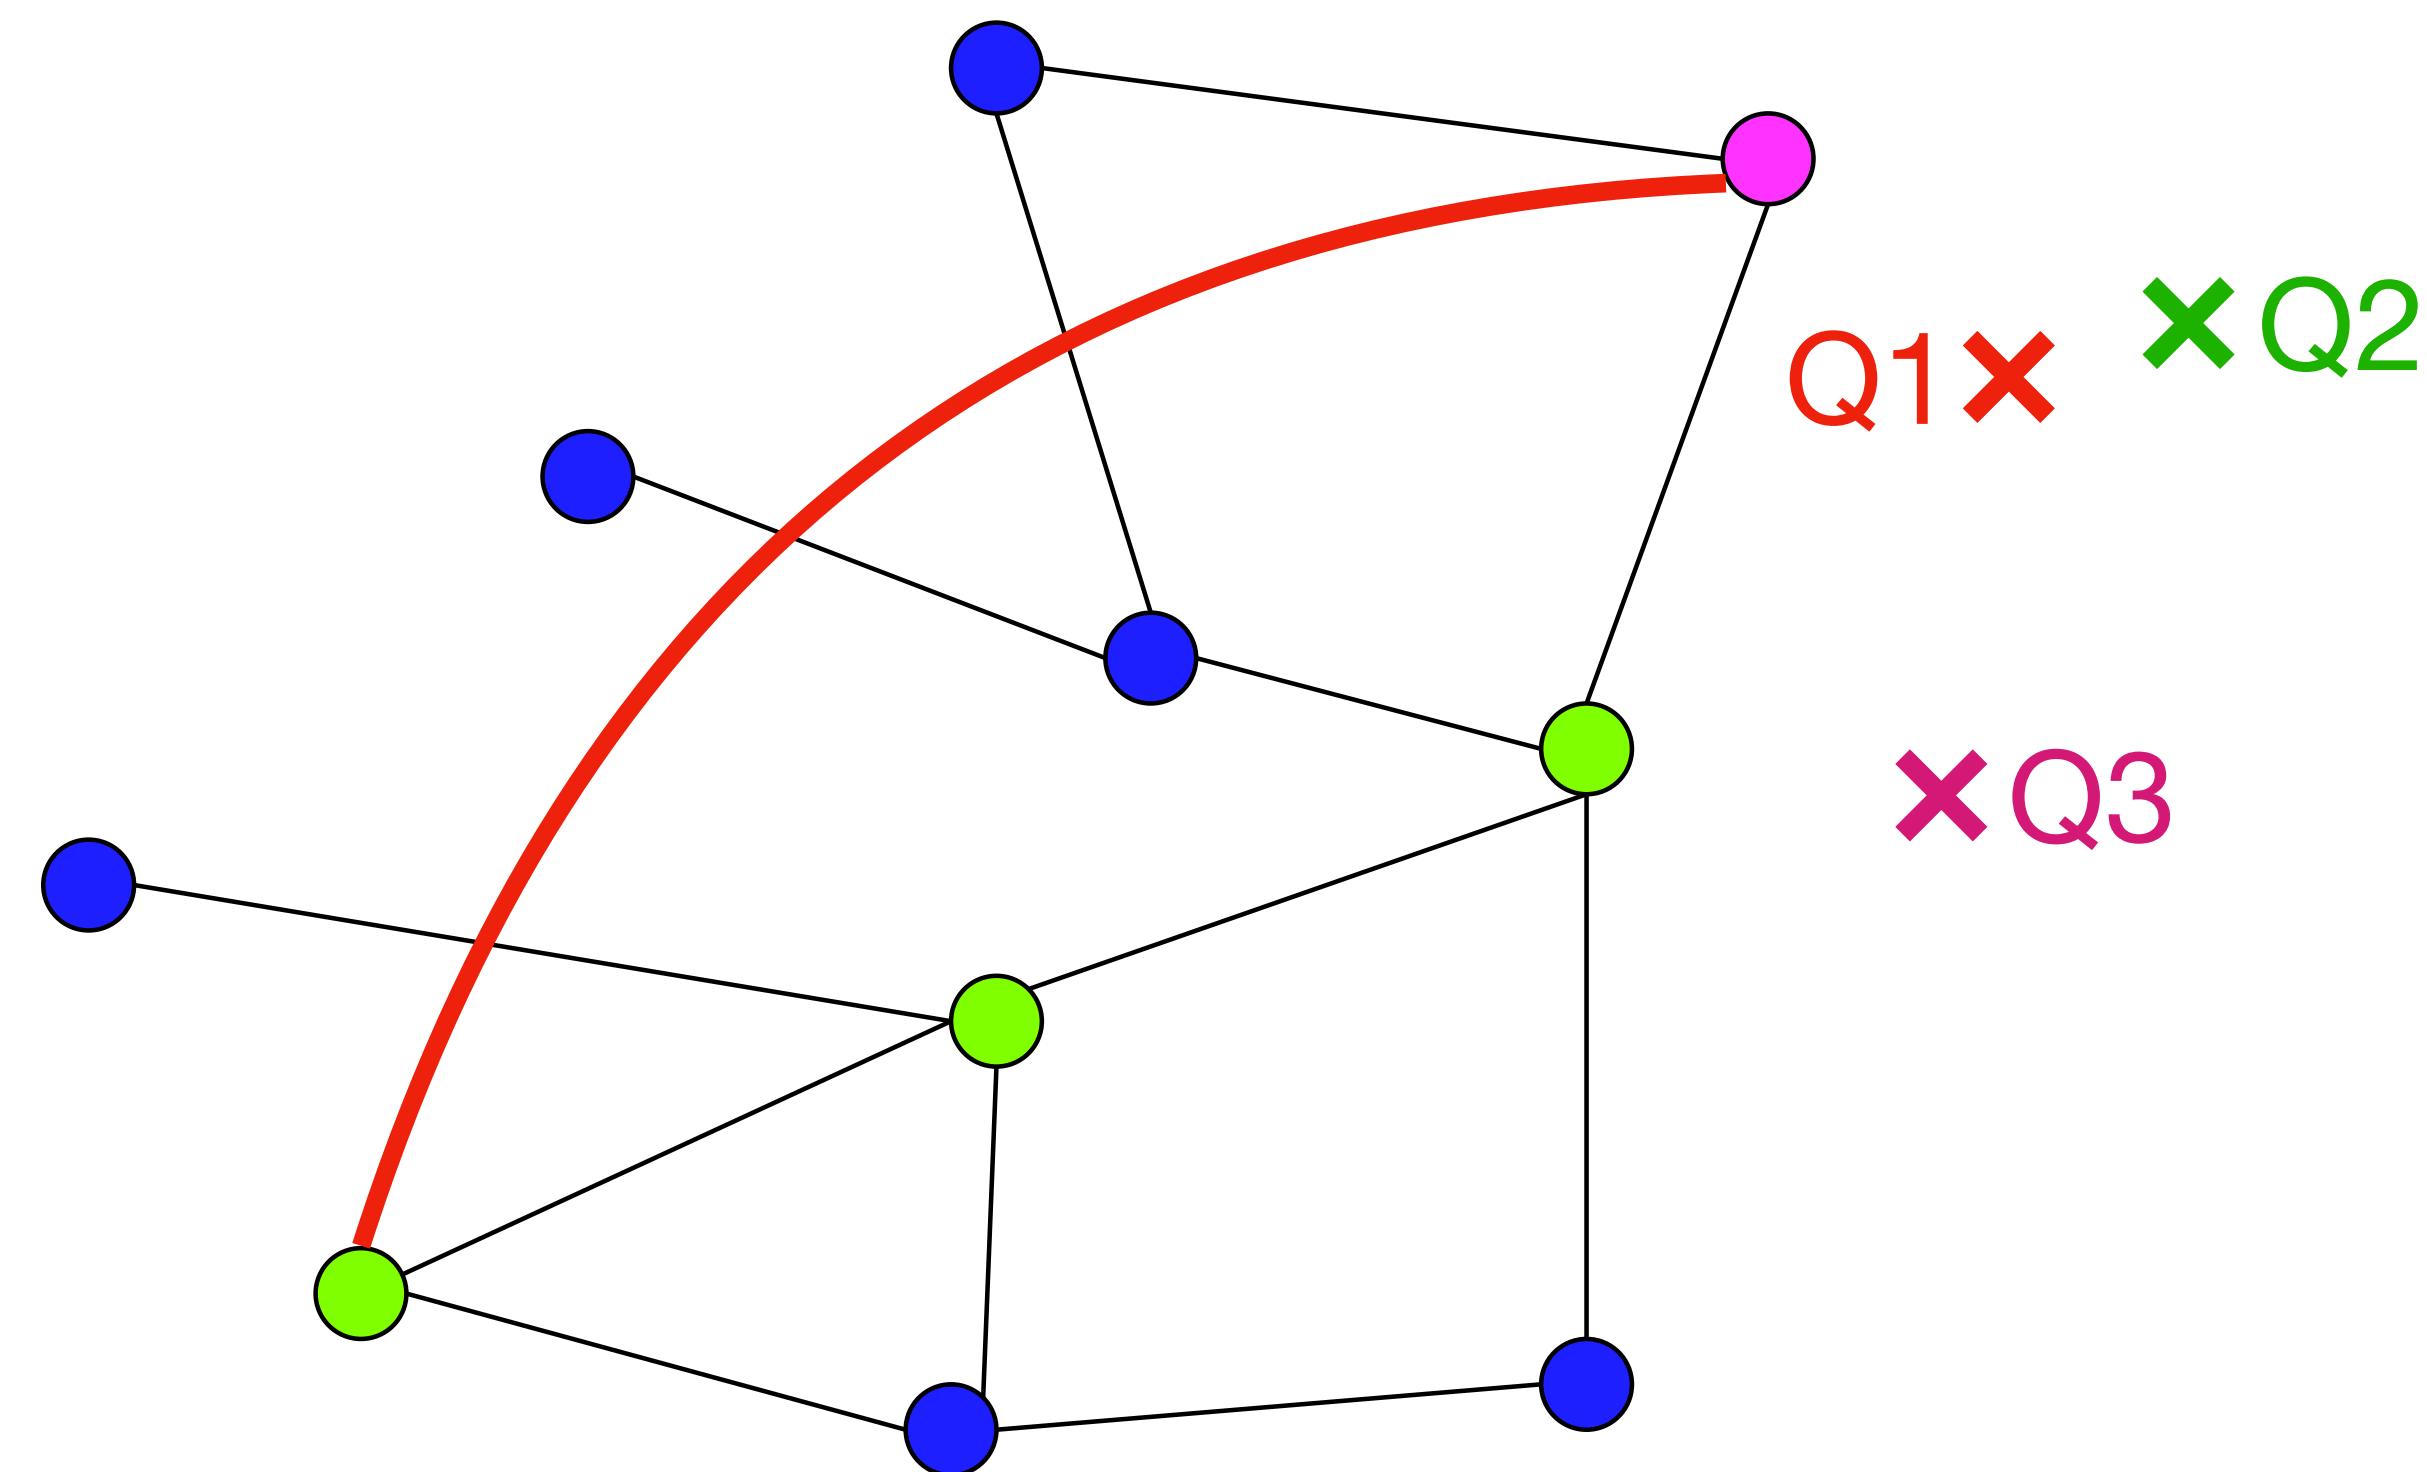
Generic graph search (baseline)

✓ implemented



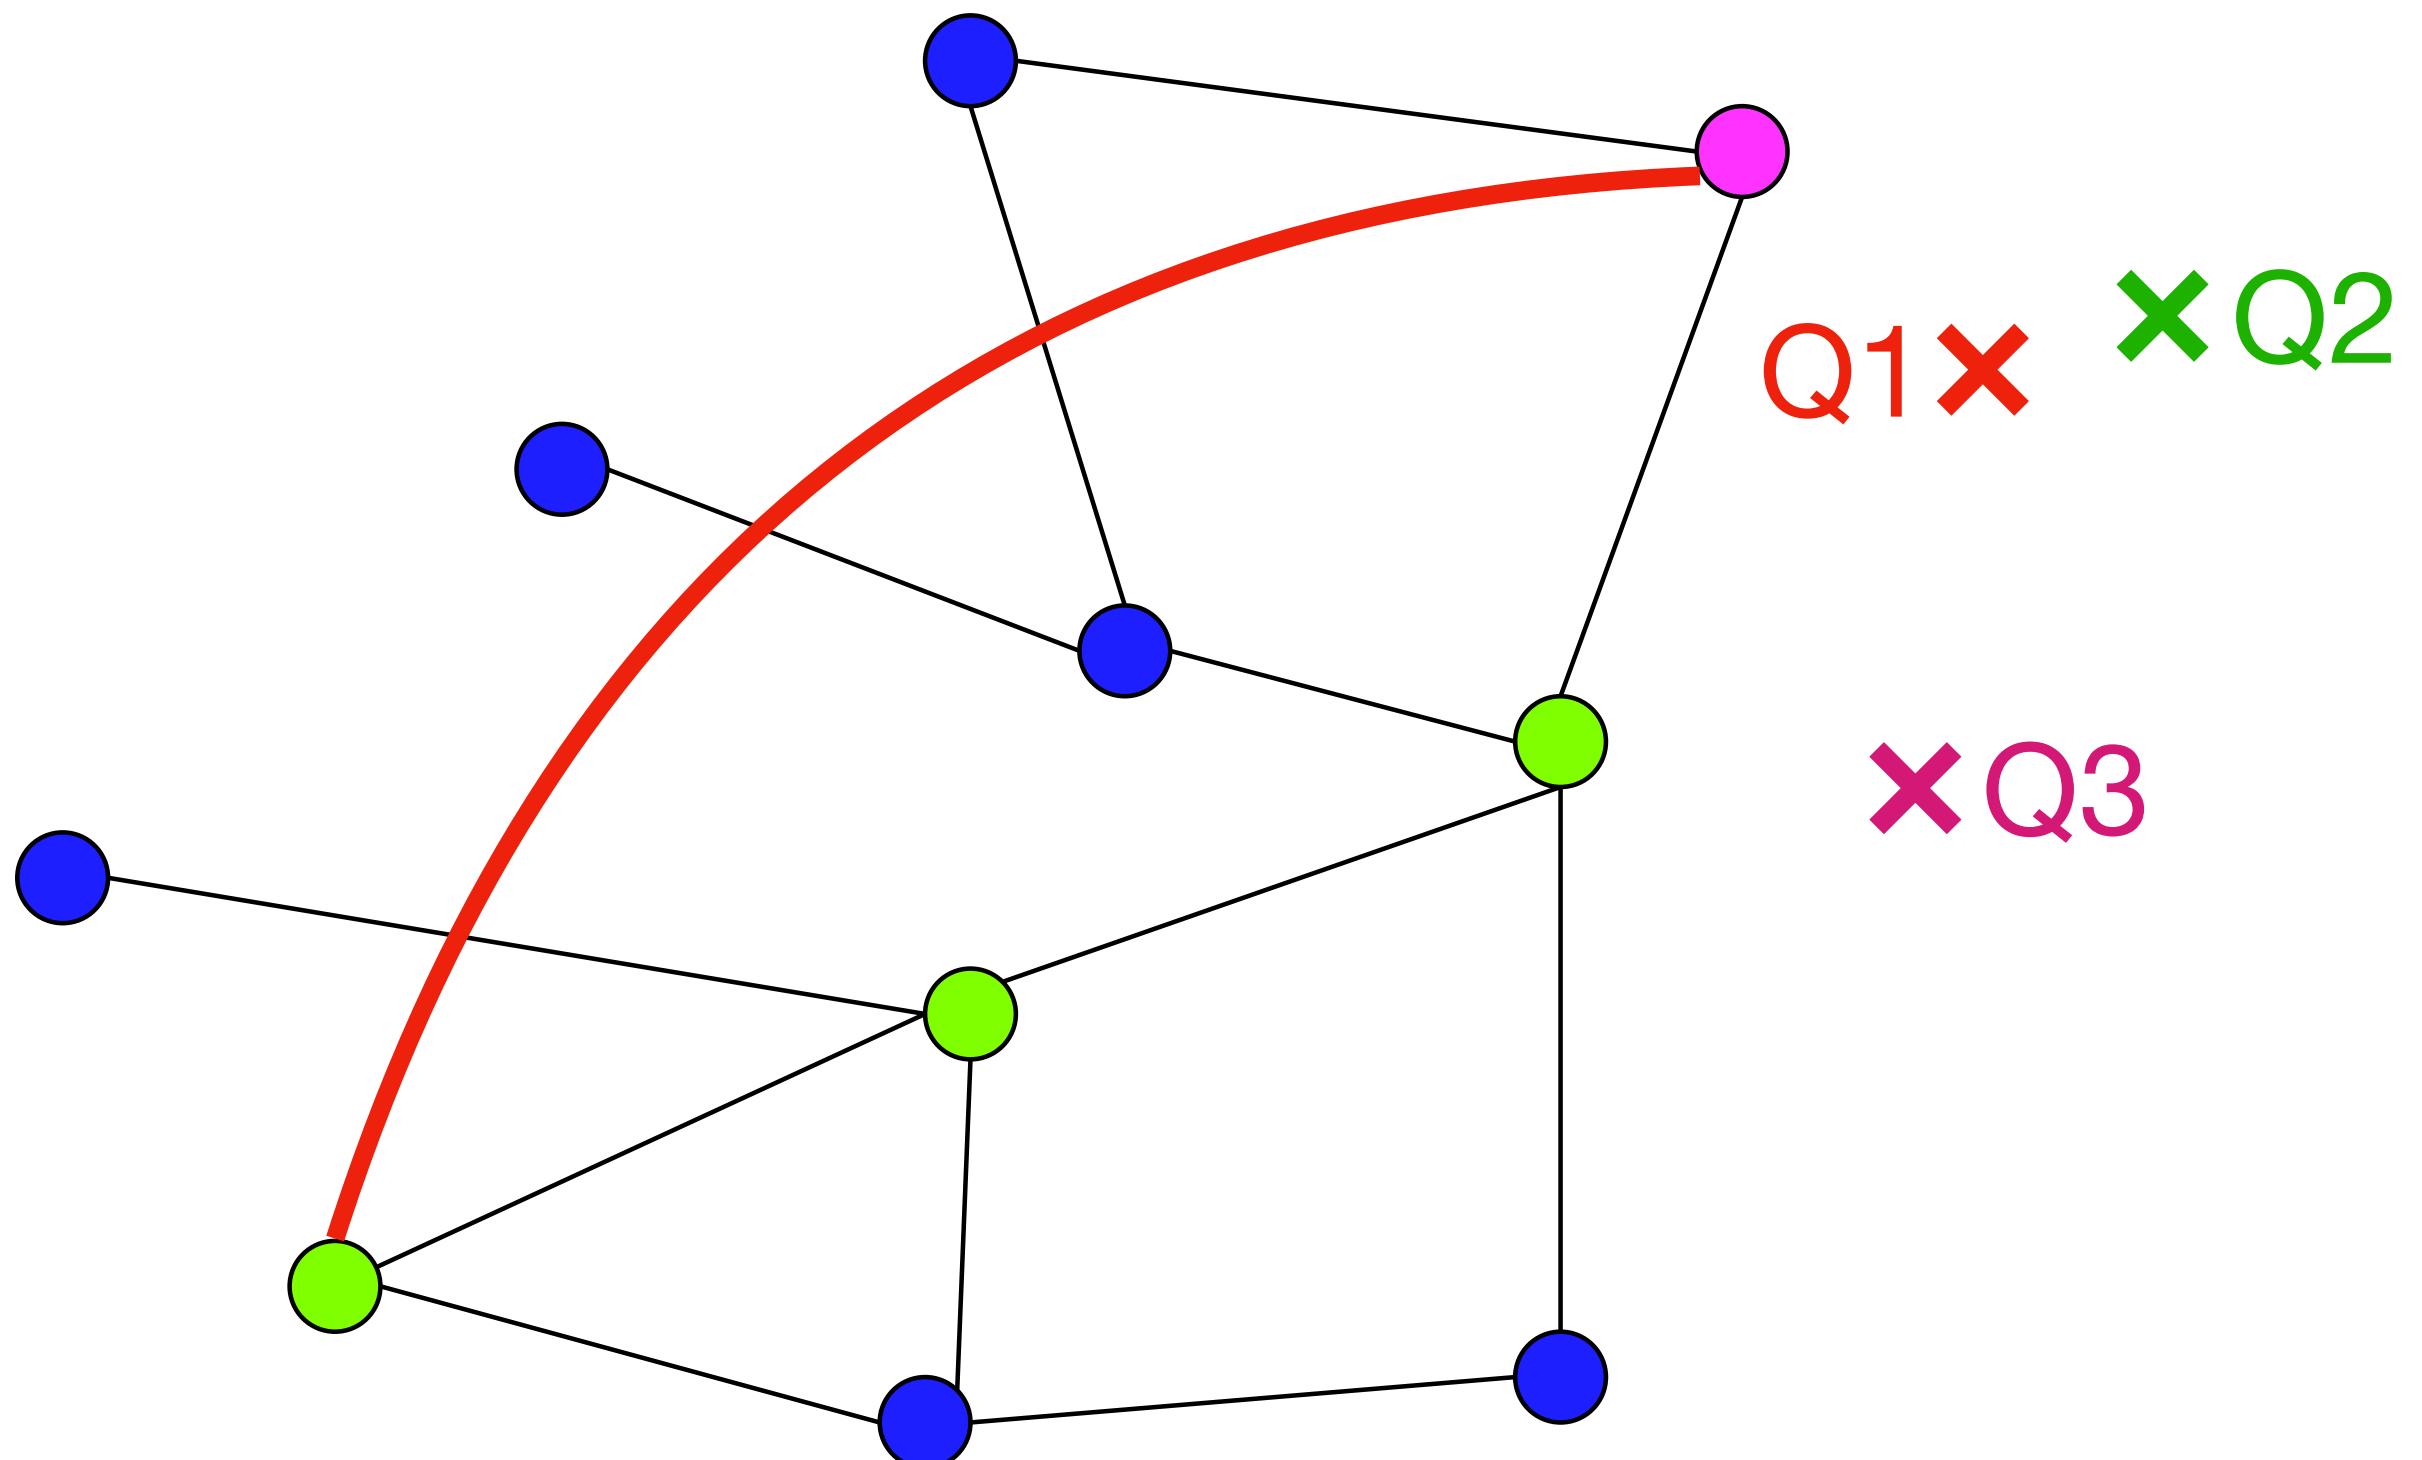
Catapult management

■ Implemented with FIFO only



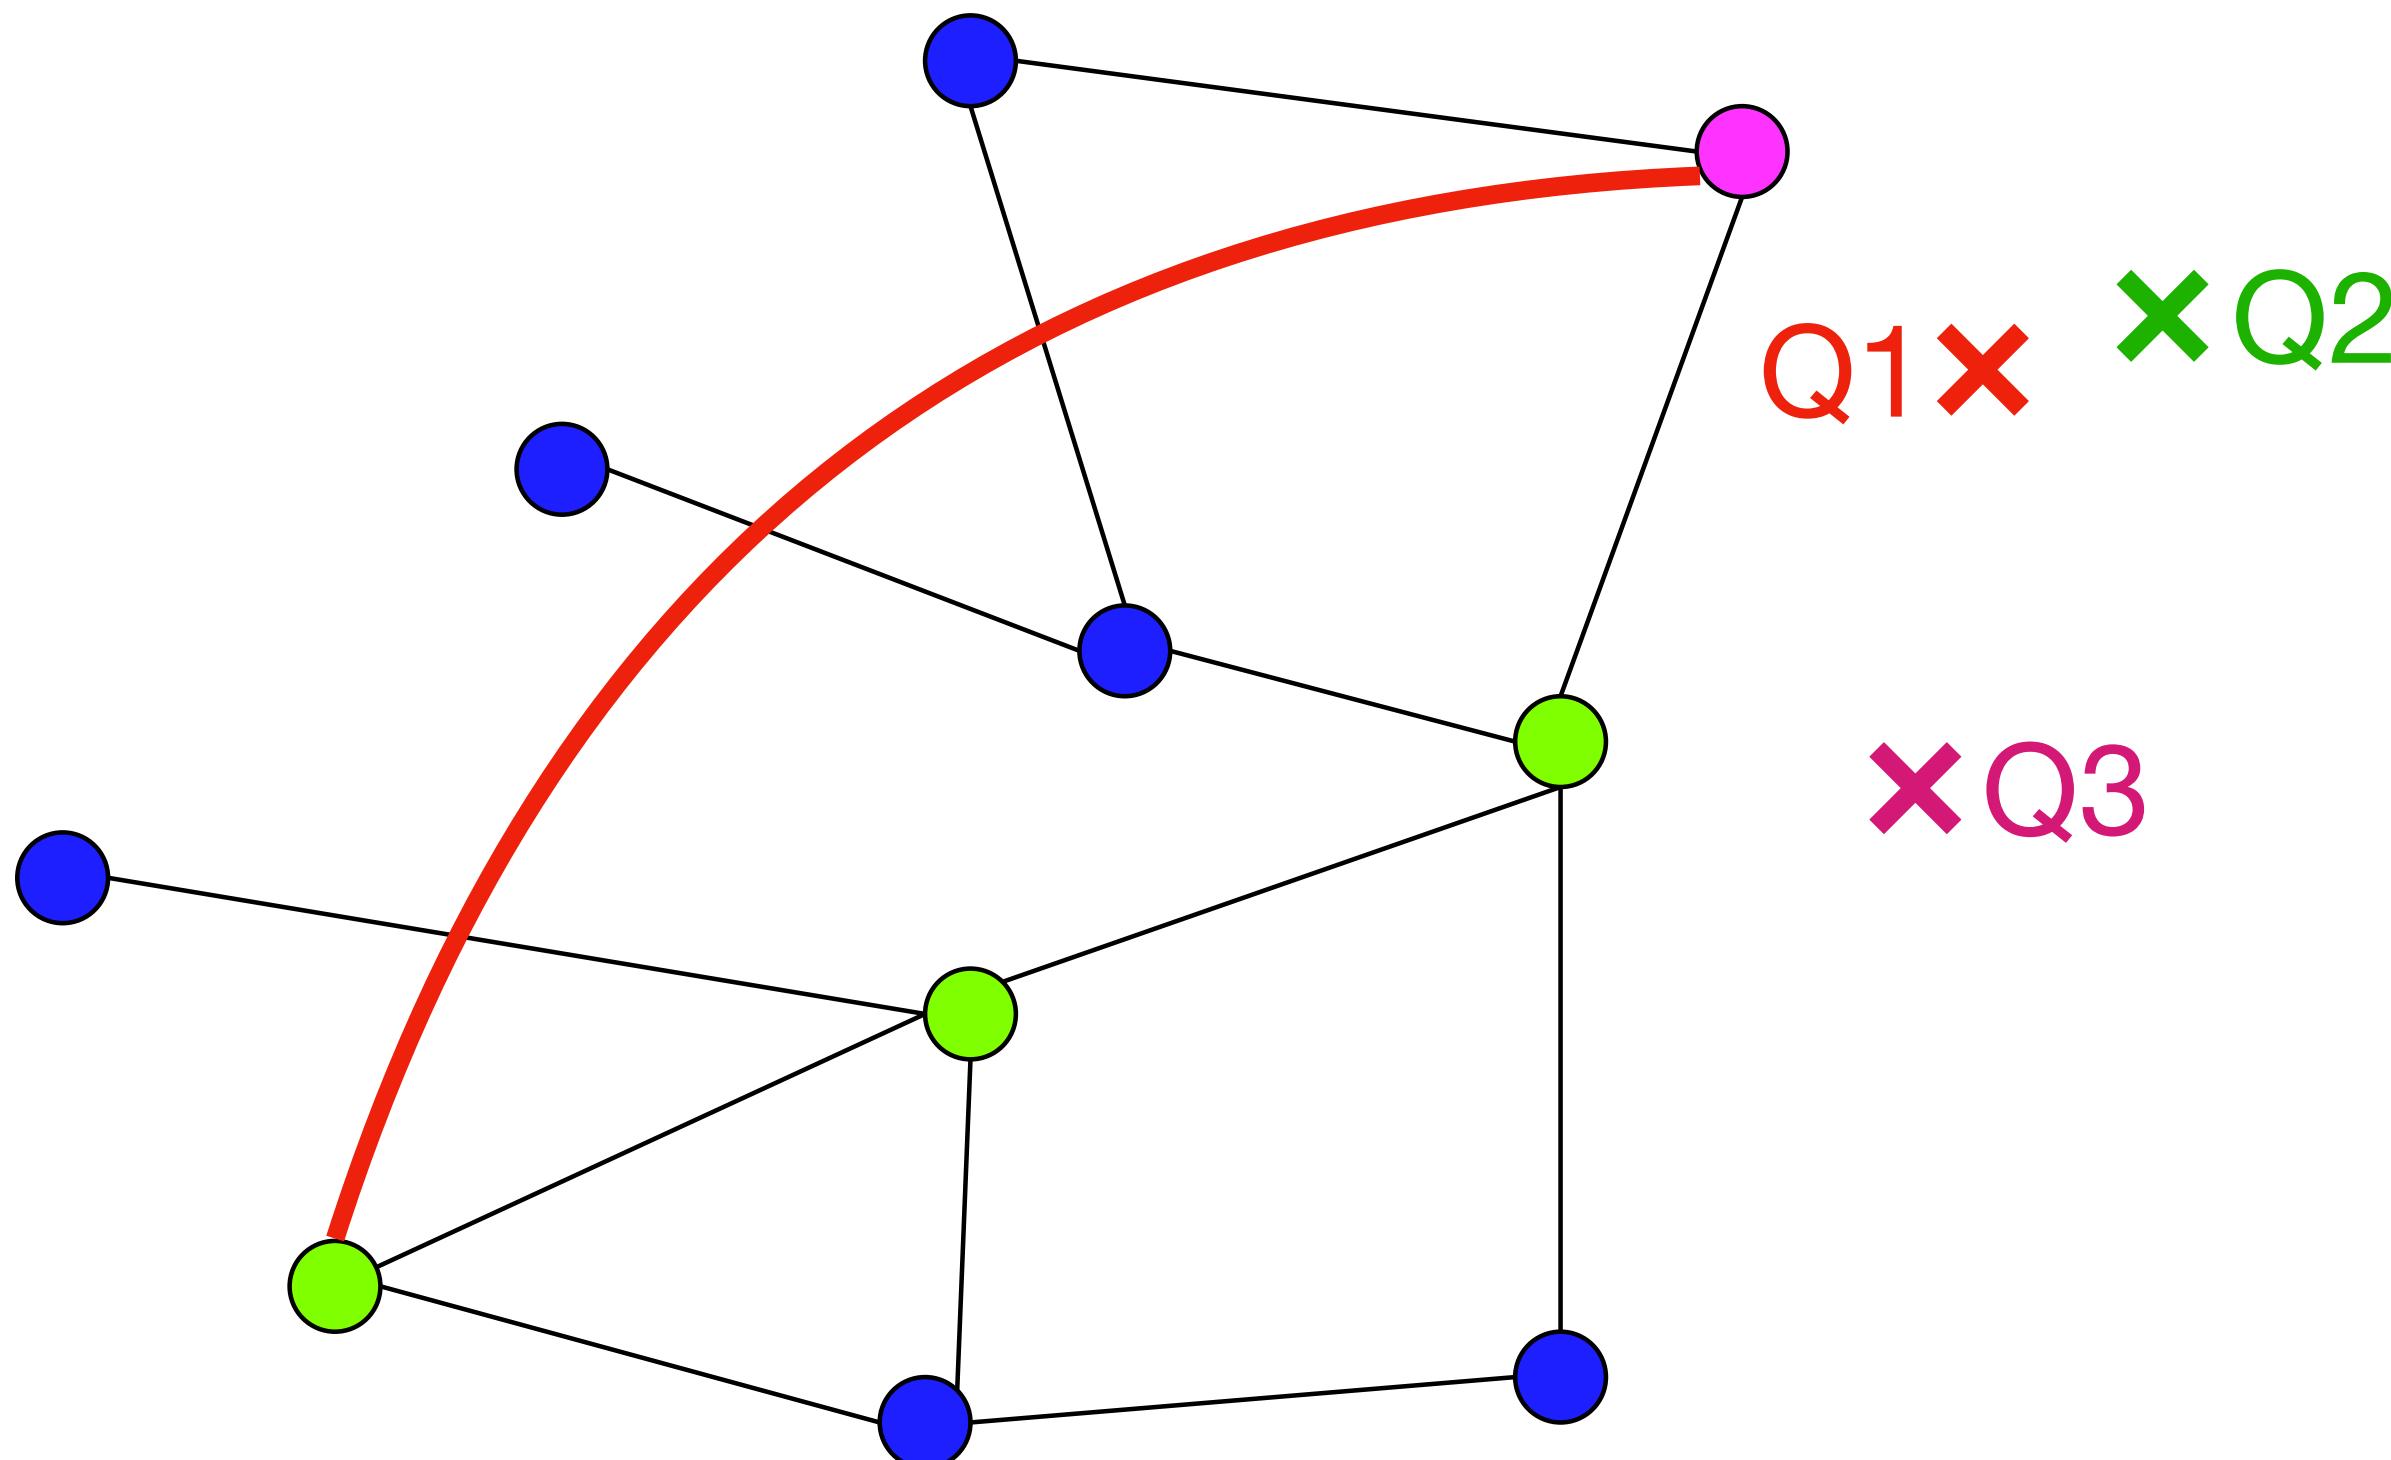
LSH for entry-point selection

✓ implemented



LSH for entry-point selection

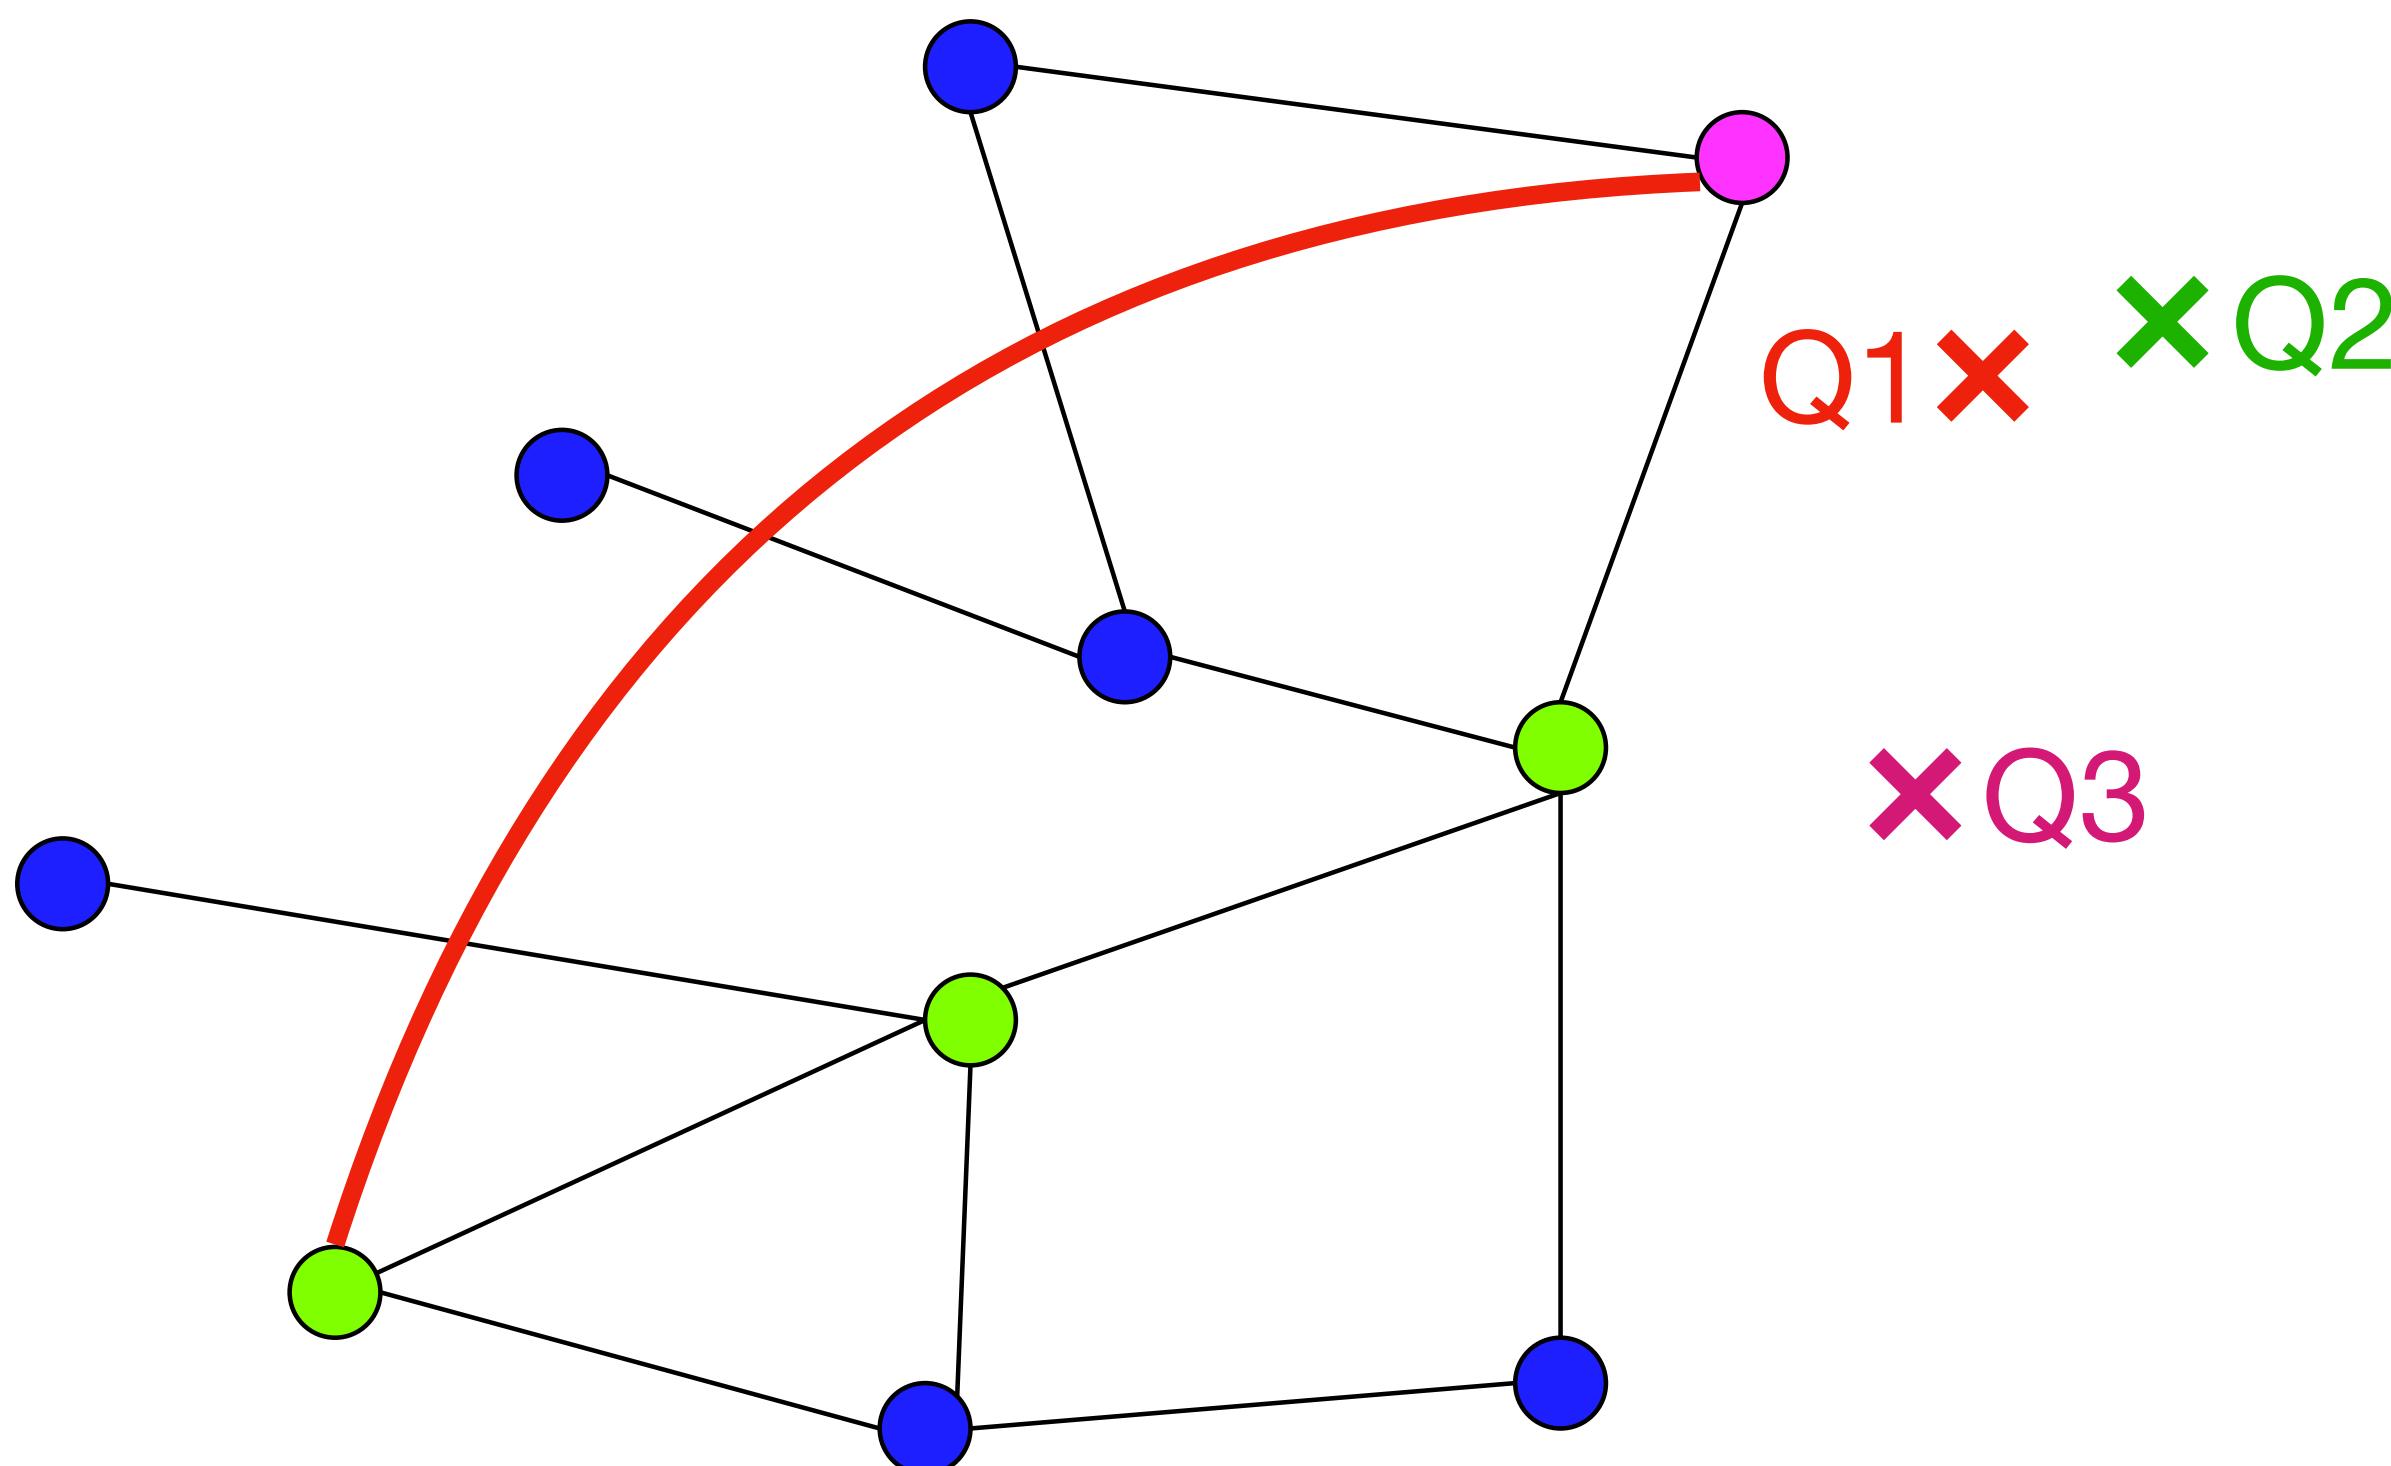
✓ implemented



Hash bits	Starting nodes
11001	1, 4, 5
10010	9, 11, 18
...	...

LSH for entry-point selection

✓ implemented



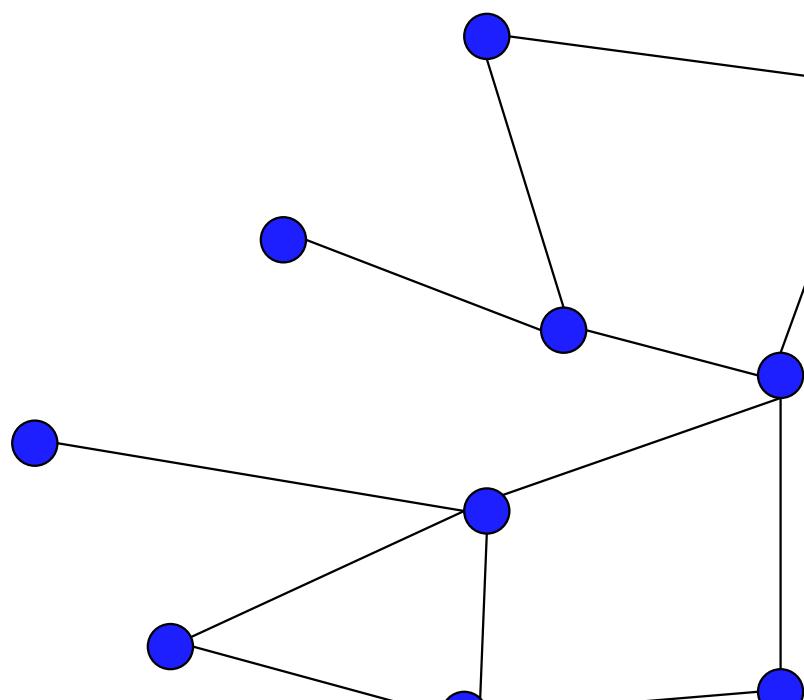
Hash bits	Starting nodes
11001	1, 4, 5
10010	9, 11, 18
...	...

New entries
are random

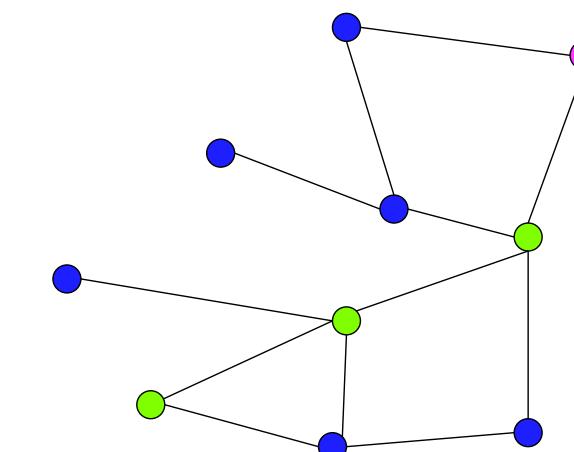
What now?

Hijacking the initial graph building

WIP



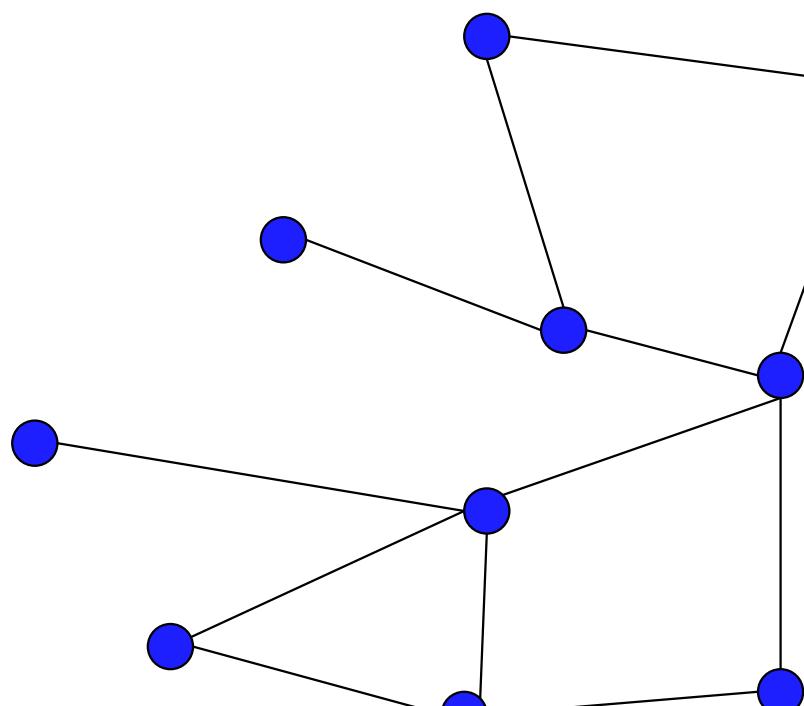
Vamana builds
a graph



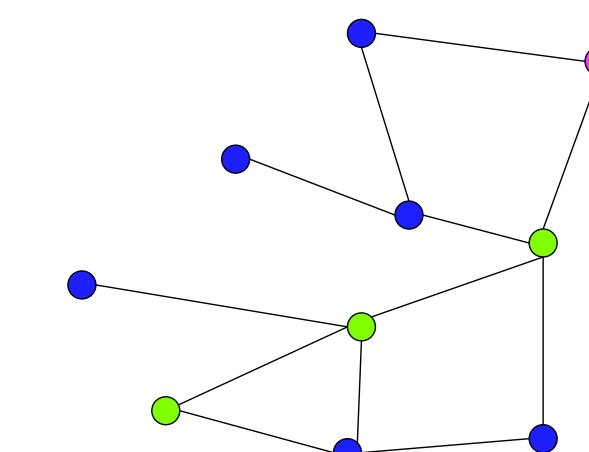
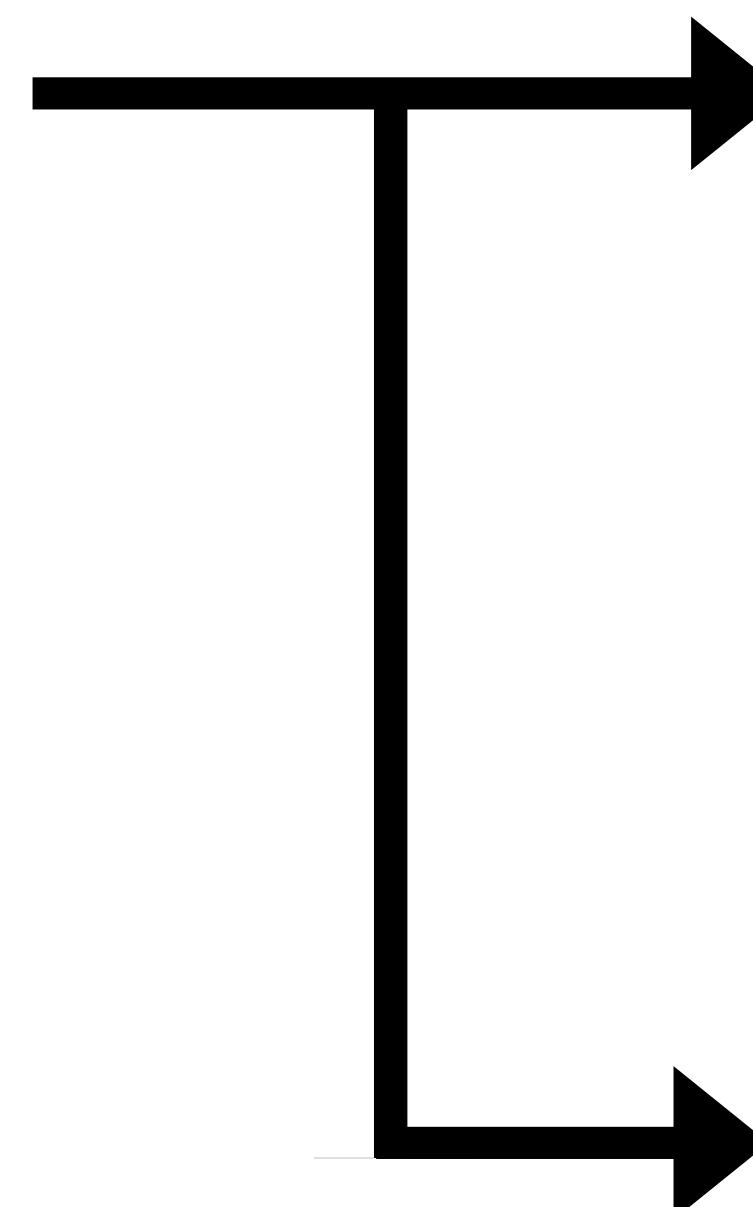
DiskANN
exploits it

Hijacking the initial graph building

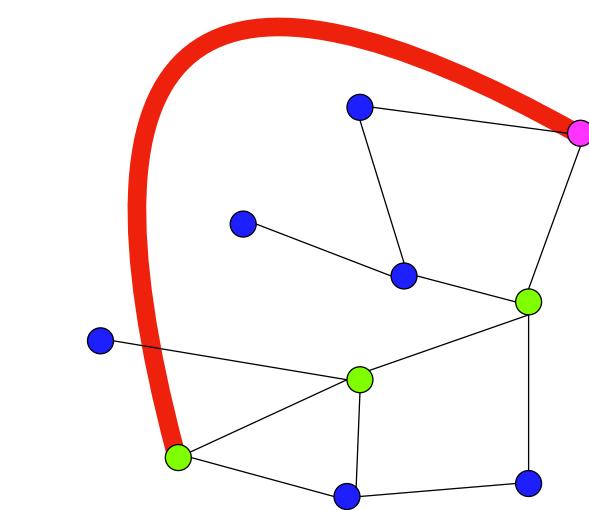
WIP



Vamana builds
a graph



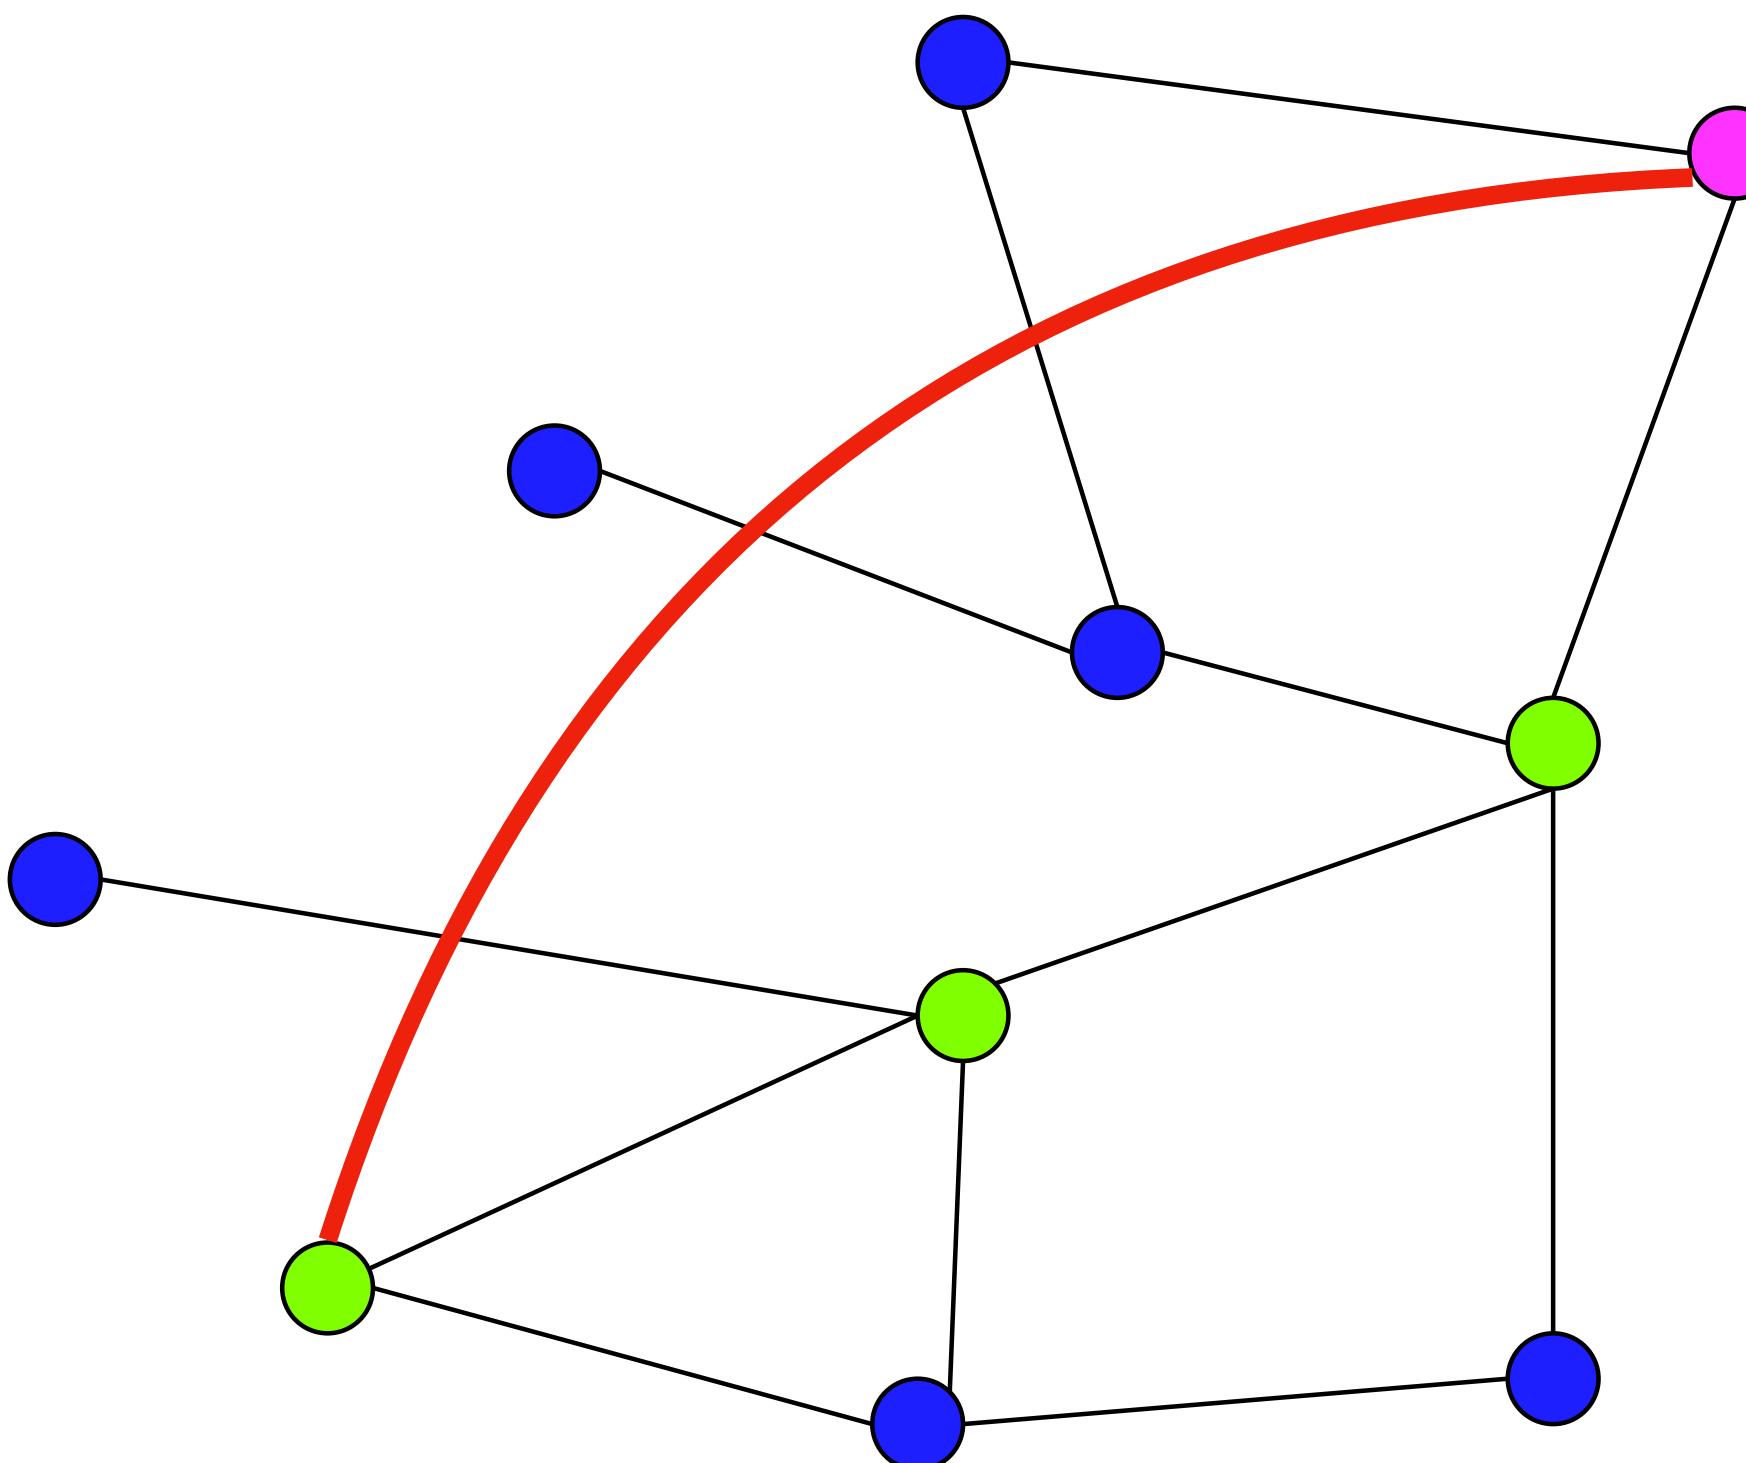
DiskANN
exploits it



Catapult
exploits it

Multithreading

✗ not implemented



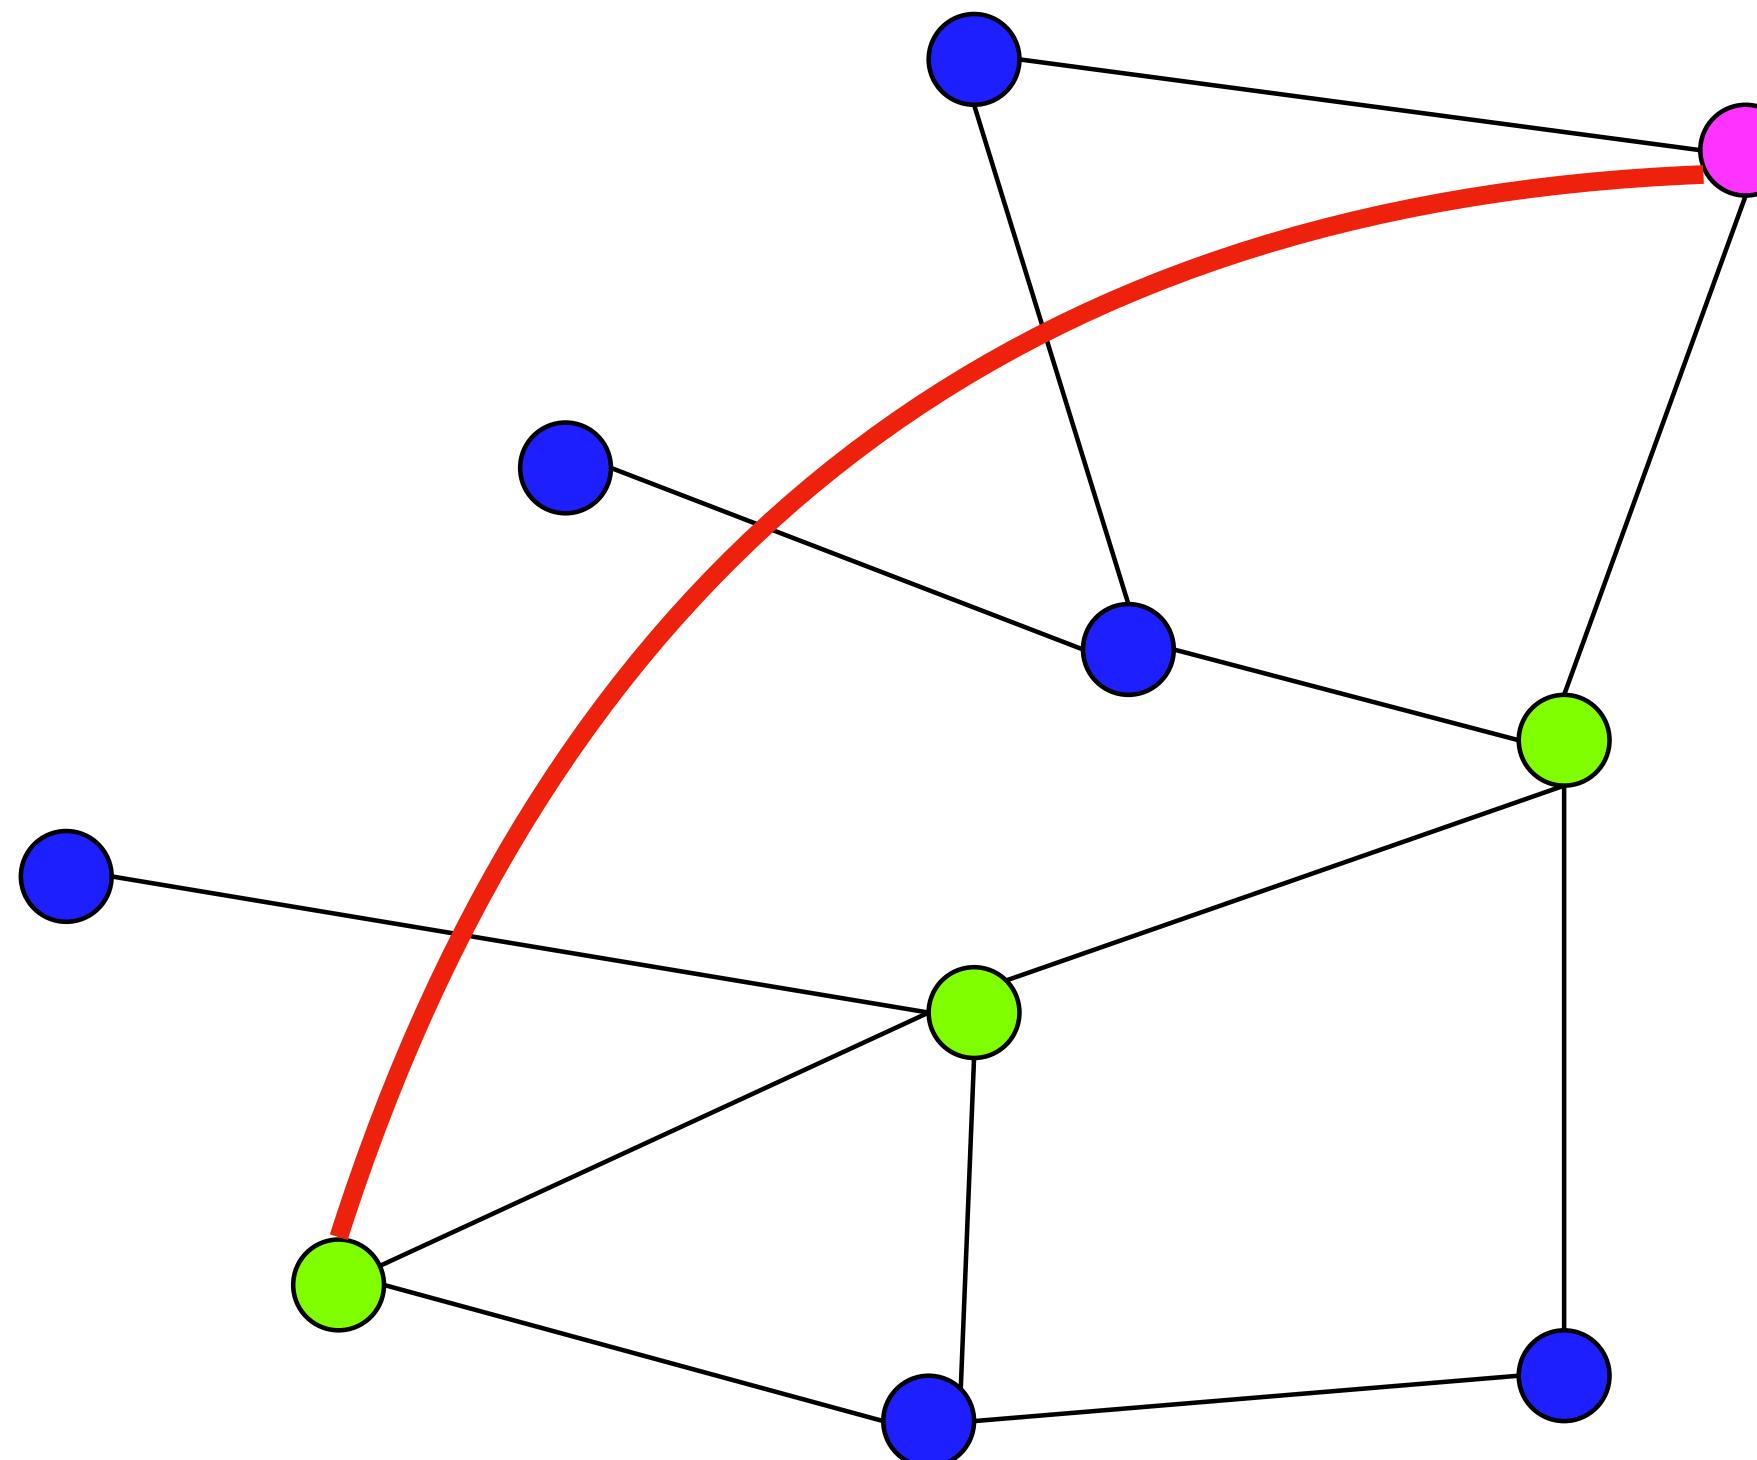
```
3  pub struct Node {  
4      pub neighbors: NeighborSet,  
5      pub catapults: NeighborSet,  
6      pub payload: Box<[f32]>,  
7  }
```

```
1  pub type NeighborSet = Vec<u32>;
```

-> Not thread-safe

Profiling / CI / regression analysis

✗ not implemented



<- What is slow in there?
What should be optimized?

Core experiment of paper

- Show we are better than baseline and/or DiskANN C++ implementation
- #hops, #edges explored, ~~wall time~~, latency, throughput, quality of vectors
- Involve some other databases? FAISS-HNSW?
- CPU only (Catapult/graph-based things not GPU friendly)

Conclusion

- « Tout va bien dans le meilleur des mondes » [1]

[1] *Candide*, Voltaire